

# Programming Using the Message Passing Paradigm

Chieh-Sen (Jason) Huang

Department of Applied Mathematics

National Sun Yat-sen University

Thank Ananth Grama, Anshul Gupta, George Karypis, and Vipin Kumar for providing slides.

# Topic Overview

- Topologies and Embedding
- Groups and Communicators
- Collective Communication and Computation Operations

# Topologies and Embeddings

- MPI allows a programmer to organize processors into logical k-dimensional meshes.
- The processor ids in `MPI_COMM_WORLD` can be mapped to other communicators (corresponding to higher-dimensional meshes) in many ways.
- The goodness of any such mapping is determined by the interaction pattern of the underlying program and the topology of the machine.
- MPI does not provide the programmer any control over these mappings.

# Topologies and Embeddings

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15


(a) Row-major mapping

0	4	8	12
1	5	9	13
2	6	10	14
3	7	11	15

(b) Column-major mapping

0	3	4	5
1	2	7	6
14	13	8	9
15	12	11	10

(c) Space-filling curve mapping



(d) 4-d Hypercube.

Different ways to map a set of processes to a two-dimensional grid. (a) and (b) show a row- and column-wise mapping of these processes, (c) shows a mapping that follows a space-filling curve (dotted line), and (d) shows a mapping in which neighboring processes are directly connected in a hypercube.

# Creating and Using Cartesian Topologies

- We can create cartesian topologies using the function:

```
int MPI_Cart_create(MPI_Comm comm_old, int ndims, int *dims,  
                   int *periods, int reorder, MPI_Comm *comm_cart)
```

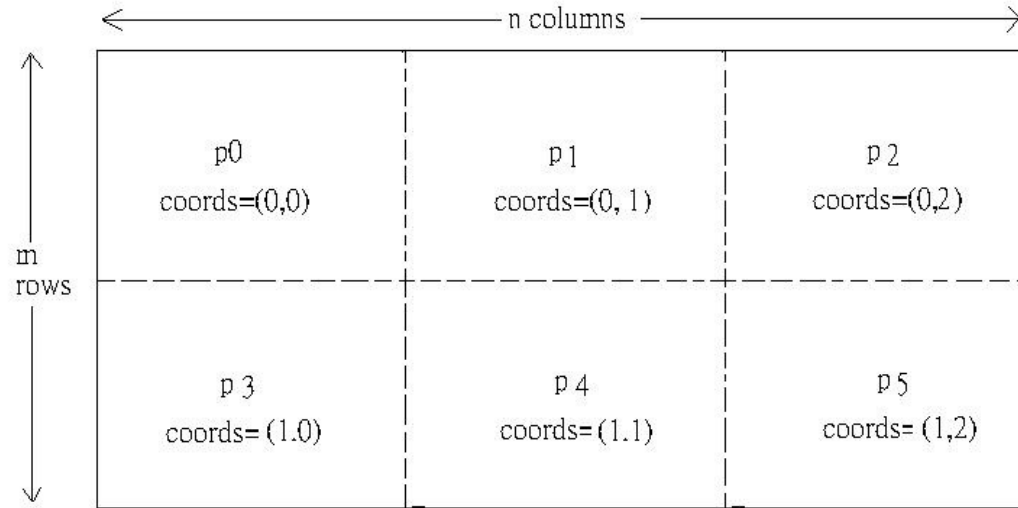
This function takes the processes in the old communicator and creates a new communicator with `dims` dimensions.

- Each processor can now be identified in this new cartesian topology by a vector of dimension `dims`.

- 
- `ndims` = number of dimensions.
  - `*dims` → the number of elements in each dimension.
  - `periods` → nonzero (True) for wrapping.
  - `reorder` → reorder the processor.

# Creating and Using Cartesian Topologies

```
dims[0]=2;  
dims[1]=3;  
periods[0]=periods[1]=1;
```



2D array distributed on a 2x3 process grid

- Since sending and receiving messages still require (one-dimensional) ranks, MPI provides routines to convert ranks to cartesian coordinates and vice-versa.

```
int MPI_Cart_coord(MPI_Comm comm_cart, int rank, int maxdims,  
                  int *coords)  
int MPI_Cart_rank(MPI_Comm comm_cart, int *coords, int *rank)
```

# Creating and Using Cartesian Topologies

- The most common operation on cartesian topologies is a shift. To determine the rank of source and destination of such shifts, MPI provides the following function:

```
int MPI_Cart_shift(MPI_Comm comm_cart, int dir, int s_step,  
                  int *rank_source, int *rank_dest)  
MPI_Cart_shift(comm_2d, 1, -1, &shiftsource, &shiftdest);  
// 1: in 2nd dim (row-wise), -1: to the left
```

- 
- `dir`: 0 (1st dim, column-wise ↓), 1 (2nd dim, row-wise →).
  - `s_step`: Positive, 1st dim, column-wise ↓, 2nd dim, row-wise →.

# Groups and Communicators

- In many parallel algorithms, communication operations need to be restricted to certain subsets of processes.
- MPI provides mechanisms for partitioning the group of processes that belong to a communicator into subgroups each corresponding to a different communicator.





# Groups and Communicators

- In many parallel algorithms, processes are arranged in a virtual grid, and in different steps of the algorithm, communication needs to be restricted to a different subset of the grid.
- MPI provides a convenient way to partition a Cartesian topology to form lower-dimensional grids:

```
int MPI_Cart_sub(MPI_Comm comm_cart, int *keep_dims,  
                MPI_Comm *comm_subcart)  
keep_dims[0]=1;keep_dims[1]=0;  
// 3 column sub-topologies are created
```

- If `keep_dims[i]` is true (non-zero value in C) then the `i`th dimension is retained in the new sub-topology.
- The coordinate of a process in a sub-topology created by `MPI_Cart_sub` can be obtained from its coordinate in the original topology by disregarding the coordinates that correspond to the dimensions that were not retained.

# Groups and Communicators




Splitting a Cartesian topology of size  $2 \times 4 \times 7$  into (a) four subgroups of size  $2 \times 1 \times 7$ , and (b) eight subgroups of size  $1 \times 1 \times 7$ .

- Homework: Hypercube example.

# Collective Communication and Computation Operations

- MPI provides an extensive set of functions for performing common collective communication operations.
- Each of these operations is defined over a group corresponding to the communicator.
- All processors in a communicator must call these operations.
- The barrier synchronization operation is performed in MPI using:

```
int MPI_Barrier(MPI_Comm comm)
```



# Compute the Global Sum

```
int *global_sum;
....
global_sum=(int*)malloc((npes)*sizeof(int));
for(i=0;i<4;i++){
    a[i]= myrank*4 + i;
    local_sum += a[i];
}
printf("My rank %d a= %d,%d,%d,%d sum= %d \n",myrank,a[0],a[1],a[2],a[3],local_sum);

if(myrank !=0)
    MPI_Send(&local_sum,1,MPI_INT,0,1,MPI_COMM_WORLD);
else
    global_sum[0]= local_sum;
if(myrank == 0){
    for(i=1;i<npes;i++)
        MPI_Recv(global_sum+i,1,MPI_INT,i,1,MPI_COMM_WORLD,&status);
}
MPI_Barrier(MPI_COMM_WORLD);
if(myrank ==0){
    for(i=1;i<npes;i++)
        global_sum[0] += global_sum[i];
    printf("My rank %d sum = %d \n",myrank, global_sum[0]);
}
MPI_Bcast(global_sum,1,MPI_INT,0,MPI_COMM_WORLD);
printf("My rank %d sum = %d \n",myrank, global_sum[0]);
....
```

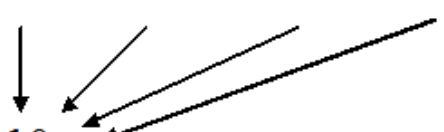
# Collective Communication Operations

- The one-to-all broadcast operation is:

```
int MPI_Bcast(void *buf, int count, MPI_Datatype datatype,
              int source, MPI_Comm comm)
```

- The all-to-one reduction operation is:

```
int MPI_Reduce(void *sendbuf, void *recvbuf, int count,
               MPI_Datatype datatype, MPI_Op op, int target,
               MPI_Comm comm)
```

MPI_Scan	$p_0$	$p_1$	$p_2$	$p_3$
	1	2	3	4
MPI_OP MPI_MAX	1	2	3	4
MPI_OP MPI_SUM	1	3	6	10
MPI_Reduce				
MPI_sum	10			
MPI_Allreduce				
MPI_sum	10	10	10	10

# Predefined Reduction Operations

Operation	Meaning	Datatypes
MPI_MAX	Maximum	C integers and floating point
MPI_MIN	Minimum	C integers and floating point
MPI_SUM	Sum	C integers and floating point
MPI_PROD	Product	C integers and floating point
MPI LAND	Logical AND	C integers
MPI_BAND	Bit-wise AND	C integers and byte
MPI_LOR	Logical OR	C integers
MPI BOR	Bit-wise OR	C integers and byte
MPI_LXOR	Logical XOR	C integers
MPI_BXOR	Bit-wise XOR	C integers and byte
MPI_MAXLOC	max-min value-location	Data-pairs
MPI_MINLOC	min-min value-location	Data-pairs

# Collective Communication Operations

MPI datatypes for data-pairs used with the `MPI_MAXLOC` and `MPI_MINLOC` reduction operations.

MPI Datatype	C Datatype
<code>MPI_2INT</code>	pair of ints
<code>MPI_SHORT_INT</code>	short and int
<code>MPI_LONG_INT</code>	long and int
<code>MPI_LONG_DOUBLE_INT</code>	long double and int
<code>MPI_FLOAT_INT</code>	float and int
<code>MPI_DOUBLE_INT</code>	double and int

# Collective Communication Operations

- If the result of the reduction operation is needed by all processes, MPI provides:

```
int MPI_Allreduce(void *sendbuf, void *recvbuf, int count,  
                 MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```

- To compute prefix-sums, MPI provides:

```
int MPI_Scan(void *sendbuf, void *recvbuf, int count,  
            MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```

- Homework: Using collective operators to write a mpi code to compute standard deviation of n numbers using m processors.





# Collective Communication Operations

- The gather operation is performed in MPI using:

```
int MPI_Gather(void *sendbuf, int sendcount,  
              MPI_Datatype senddatatype, void *recvbuf, int recvcount,  
              MPI_Datatype recvdatatype, int target, MPI_Comm comm)
```

- MPI also provides the MPI\_Allgather function in which the data are gathered at all the processes.

```
int MPI_Allgather(void *sendbuf, int sendcount,  
                 MPI_Datatype senddatatype, void *recvbuf, int recvcount,  
                 MPI_Datatype recvdatatype, MPI_Comm comm)
```



# Collective Communication Operations



The corresponding scatter operation is:

```
int MPI_Scatter(void *sendbuf, int sendcount,  
               MPI_Datatype senddatatype, void *recvbuf, int recvcount,  
               MPI_Datatype recvdatatype, int source, MPI_Comm comm)
```


- The all-to-all personalized communication operation is performed by:

```
int MPI_Alltoall(void *sendbuf, int sendcount,  
                 MPI_Datatype senddatatype, void *recvbuf, int recvcount,  
                 MPI_Datatype recvdatatype, MPI_Comm comm)
```


- Using this core set of collective operations, a number of programs can be greatly simplified.




# Collective Communication Operations



(a) Reduce( $x, y, P2, \text{MAX}$ )



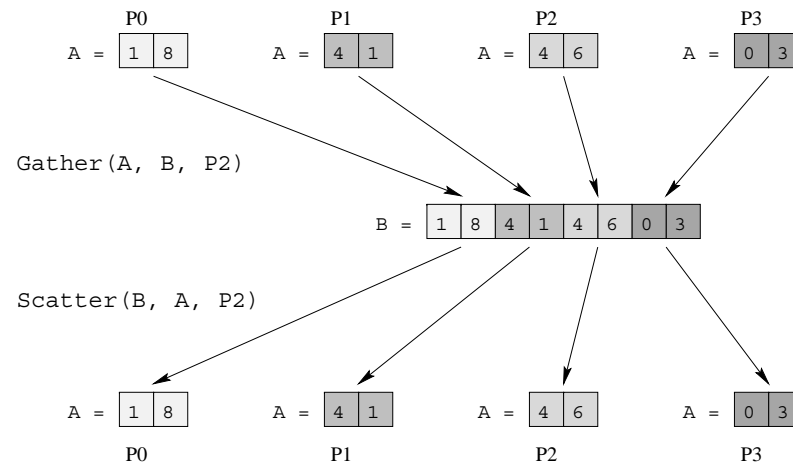
(b) Reduce( $A, B, P1, \text{SUM}$ )



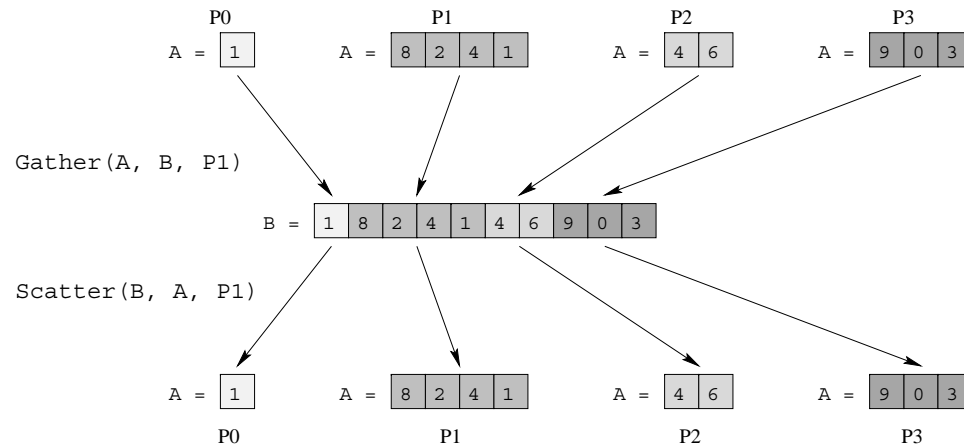
(c) AllReduce( $A, B, \text{SUM}$ )

# Collective Communication Operations


## (a) Equal-Size Gather and Scatter Operations



## (b) Unequal-Size Gather and Scatter Operations



# Collective Communication Operations



AllToAll Operation of Equal Size Data