

Remote Memory Operations

Chieh-Sen (Jason) Huang

Department of Applied Mathematics

National Sun Yat-sen University


Materials are taken from the book, Using MPI-2: Advanced Features of the Message-Passing Interface by William Gropp, Ewing Lusk, and Rajeev Thakur.

Topic Overview

- Introduction to Remote Memory Operations
- A Non-Remote Memory Operation code
- Remote Memory Operations

Introduction to Remote Memory Operations

- Send-Recv:
Specify what memory locations can be modified.
Specify when this can happen.
- Advantages:
 1. Ensuring the correctness of the data.
 2. Changes in memory of the receiver can happen only when the receiver allows.
- Disadvantages:
 1. Expressiveness and low performance.
 2. What if data that is updated frequently.
 3. Communication costs might be big.



Remote Memory Access (RMA)

- One single routine that specifies both where the data is coming from and where it is going to.
- They are called one-sided communication routines.
- Three main steps when using the RMA.
 1. Define the memory of the process that can be used for RMA operations.
`MPI_Win_create`
 2. Specify the data to be moved and where to move it.
`MPI_Put, MPI_Get, MPI_Accumulate`
 3. Specify how we know that the data is available.
`MPI_Win_fence`

A Non-RMA code

$$\pi = \int_0^1 \frac{4}{1+x^2} dx.$$

- Compute the Riemann sum.

```
h = 1.0 / (double) n;
/* n is the number of partition over [0,1] */
sum = 0.0; /* local sum */
for (i = myid + 1; i <= n; i += numprocs) {
    x = h * ((double)i - 0.5);
    sum += (4.0 / (1.0 + x*x));
}
mypi = h * sum; /* partial Riemann sum */
/* reduce all partial_sum to cpu0 */
MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0,
           MPI_COMM_WORLD);
if (myid == 0)
    printf("pi is approximately %.16f, Error is %.16f\n",
           pi, fabs(pi - PI25DT));
```

Remote Memory Access (RMA)

```
MPI_Win win;  
int MPI_Win_create(void *base, MPI_Aint size, int disp_unit,  
                  MPI_Info info, MPI_Comm comm, MPI_Win *win)  
int MPI_Win_free(MPI_Win *win)
```

- `*base`: Initial address of window.
 - `size`: Size of window in bytes.
 - `disp_unit`: Local unit size for displacements, in bytes.
Unit size in different processors maybe different.
 - `info`: Info argument.
 - `*win`: Window object returned by the call.
-
- Need to be called by all processors in `comm` which contribute to this window.
 - A “window” refers to a region of memory within a single process.
 - A “window object” describes the collection of windows that are the input to the `MPI_Win_create` call.

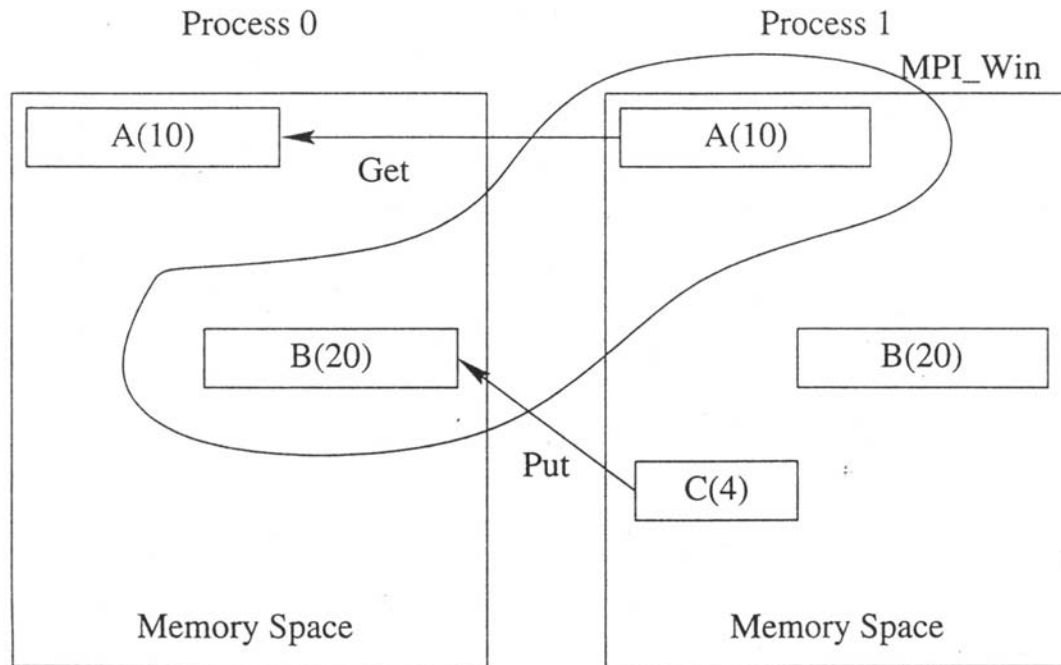
Remote Memory Access (RMA)

```
if (myid == 0){
    MPI_Win_create(&n, sizeof(int), 1, MPI_INFO_NULL,
                  MPI_COMM_WORLD, &nwin);
    MPI_Win_create(&pi, sizeof(double), 1, MPI_INFO_NULL,
                  MPI_COMM_WORLD, &piwin);
}
else{
    MPI_Win_create(MPI_BOTTOM, 0, 1, MPI_INFO_NULL,
                  MPI_COMM_WORLD, &nwin);
    MPI_Win_create(MPI_BOTTOM, 0, 1, MPI_INFO_NULL,
                  MPI_COMM_WORLD, &piwin);
}
```

- nwin consists one integer from cpu0, none from the rest.
- piwin consists one double from cpu0, none from the rest.

Remote Memory Access (RMA)–MPI_Win_create

```
if (myid == 0){  
    MPI_Win_create(B, 20*sizeof(int), sizeof(int), MPI_INFO_NULL,  
        MPI_COMM_WORLD, &win);}  
if (myid == 1){  
    MPI_Win_create(A, 10*sizeof(int), sizeof(int), MPI_INFO_NULL,  
        MPI_COMM_WORLD, &win);}
```

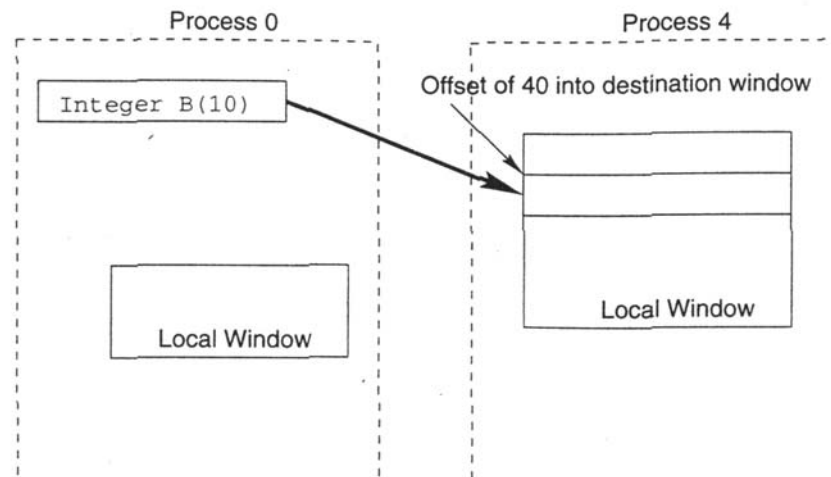


Remote Memory Access (RMA)–MPI_Put

```
int MPI_Put(void *origin_addr, int origin_count, MPI_Datatype  
            origin_datatype, int target_rank, MPI_Aint target_disp,  
            int target_count, MPI_Datatype target_datatype, MPI_Win  
            win)
```

```
-----  
if (myrank == 0) then  
    MPI_Put(B, 10, MPI_INT, 4, 40, 10, MPI_INT, win);  
-----
```

- `*origin_addr`: Initial address of origin buffer.
- `origin_count`: Number of entries in origin buffer.
- `target_rank`: Rank of target.
- `target_disp`: Displacement from window start to the beginning of the target buffer.
- `target_count`: Number of entries in target buffer.



Remote Memory Access (RMA)–MPI_Get

```
int MPI_Get(void *origin_addr, int origin_count, MPI_Datatype  
            origin_datatype, int target_rank, MPI_Aint target_disp,  
            int target_count, MPI_Datatype target_datatype, MPI_Win  
            win)
```

- `*origin_addr`: Address of the buffer in which to receive the data.
- `origin_count`: Number of entries in origin buffer.
- `target_rank`: Rank of target.
- `target_disp`: Displacement from window start to the beginning of the target buffer.
- `target_count`: Number of entries in target buffer.

Remote Memory Access (RMA)–MPI_Win_fence

- What if two processors try to access the same data at the same time?
- Answer: Use MPI_Win_fence.

```
int MPI_Win_fence(int assert, MPI_Win win)
```

- `assert`: Program assertion. Provide information about the fence that can be used by some MPI implementation to provide better performance.
- `win`: Window object.
- MPI_Win_fence completes any RMA operations that started since the last call to MPI_Win_fence.
- MPI_Win_fence is collective over all processes in the group associated with the window object.

Remote Memory Access (RMA)

- The easy rules:
 1. Do not overlap accesses on window.
 2. Separate non-RMA access from RMA access with MPI_Win_fence.
- It is possible to have several MPI window objects whose local windows overlap.
- BUT, because of the performance issues and the complexities of the rules for correct use, we recommend avoiding the use of overlapping window.
- MPI_Accumulate.

```
int MPI_Accumulate(const void *origin_addr, int origin_count,  
                  MPI_Datatype origin_datatype, int target_rank, MPI_Aint  
                  target_disp, int target_count, MPI_Datatype  
                  target_datatype, MPI_Op op, MPI_Win win)
```

Remote Memory Access (RMA)

```
if (myid == 0){
    printf("Enter the number of intervals: (0 quits) ");
    printf("\n");
    scanf("%d",&n);
    pi = 0.0;
}
MPI_Win_fence(0, nwin);

if (myid != 0)
    MPI_Get(&n, 1, MPI_INT, 0, 0, 1, MPI_INT, nwin);
MPI_Win_fence(0, nwin);

if (n == 0)
    break;
else{
    h = 1.0 / (double) n;
    sum = 0.0;
    for (i = myid + 1; i <= n; i += numprocs){
        x = h * ((double)i - 0.5);
        sum += (4.0 / (1.0 + x*x));
    }

    mypi = h * sum;

    MPI_Win_fence( 0, piwin);
    MPI_Accumulate(&mypi, 1, MPI_DOUBLE, 0, 0, 1, MPI_DOUBLE,
                  MPI_SUM, piwin);
    MPI_Win_fence(0, piwin);
}
```

Remote Memory Access (RMA)–Example

```
MPI_Win_create( A, ..., &win);
MPI_Win_fence( 0, win);

if (rank == 0)
* Process 0 puts data into many local windows */
    MPI_Put( ... , win );
/* This fence completes the MPI_Put operations initiated
   by process 0 */
MPI_Win_fence( 0, win );
/* All processes initiate access to some window to extract data */

MPI_Get( ... , win );
/* The following fence completes the MPI_Get operations */
MPI_Win_fence( 0, win );

/* After the fence, processes can load and store
   into A, the local window */
A[rank] = 4;
printf( "A[%d] = %d\n", 0, A[0] );
MPI_Win_fence( 0, win );

/* We need a fence between stores and RMA operations */
MPI_Put( ... , win );
/* The following fence completes the preceding Put */
MPI_Win_fence( 0, win );
```

Remote Memory Access (RMA)

- Homework: Use RMA to compute variance of n numbers.
- Homework: Use RMA to implement the Jacobi method for solving linear system.